



Politecnico
di Torino



Reinforcement Learning Project

April 8th, 2026

TAs: Davide Buoso (davide.buoso@polito.it), Andrea Protopapa (andrea.protopapa@polito.it)

0 Introduction

The main goal of this project is to get familiar with the reinforcement learning paradigm in the context of robotic systems. It also aims to dive into the problem of learning a control policy for a robot in simulation using state-of-the-art RL algorithms, while introducing the challenges of using that policy in the real world. In particular, the student will learn about the *sim-to-real transfer* problem in robot learning literature, namely the task of learning policies in simulation through RL that can be directly transferred to real-world hardware, avoiding costly interactions with real setups and speeding up the training time. During this project, the student will be simulating the sim-to-real transfer task in a *sim-to-sim* scenario, where a discrepancy between *source* (training) and *target* (test) domains is manually injected. The student will implement *domain randomization* of parameters (e.g. masses), a popular strategy to learn robust policies that transfer well to the target domain.

Specifically, the student must go through the following steps:

1. **Preliminaries:** Read the provided material to get familiar with the Reinforcement Learning framework, the sim-to-real transfer challenge and the common techniques to perform an efficient transfer from simulation to reality;
2. **Train your first RL agent:** Implement two basic Reinforcement Learning algorithms in literature (i.e., REINFORCE and Actor-Critic Policy Gradient algorithms) to train your first RL agent;
3. **Lower/upper bound baselines:** Implement an RL training pipeline via third-party APIs to state-of-the-art reinforcement learning algorithms such as PPO and SAC, providing the baselines for the subsequent phase;
4. **Domain Randomization:** Implement Uniform Domain Randomization (UDR) and ADR (Automatic Domain Randomization), to learn robust policies in the source domain and limit the loss of performance during the sim-to-real transfer;

For each of the aforementioned steps, a **soft deadline** is provided to guide the student along a potential schedule of the work throughout the project. This timeline is not to be seen as mandatory, but only as a suggestion of a possible weekly workload to ensure timely completion of the project.

Guiding questions, hints, and external references are also included to assist the student in comprehending the various topics independently.

Your submission will consist of (1) a **PDF report** (with paper-style) which carefully and systematically describes the entire work carried out during the project, and (2) the **code** implementation. In particular, the

report should contain a brief introduction, a related work section, a methodological section, an experimental section with all the results and discussions, and a final brief conclusion. Follow this [link](#) to open and create the template for the report using Overleaf.

In the [starting code repository](#), you will find two folders:

- **part1**: refers to steps 1 and 2
- **part2**: refers to steps 3 and 4

1 Preliminaries

Soft deadline: 15/04/2026

Before starting, you're asked to take some time to familiarize yourself with the framework of Reinforcement Learning, the sim-to-real transfer challenge and SOTA strategies to overcome it. More in detail:

- Read sections 1.-1.4, 1.6, 3.-3.8, of [1] to understand the general Reinforcement Learning framework;
- Watch introductory video on Reinforcement learning by DeepMind [video](#)
- Read [article](#) on introduction to Reinforcement learning by OpenAI [part 1, part 2, part 3]
- Read sections 1.-1.3, 3.-3.4, of [2]
- Read sections 1., 2., 3. of [3]
- Read debate on the sim-to-real transfer paradigm [4]
- Read [5], [6], [9], [blog post](#) to understand domain randomization for sim-to-real transfer
- Read this [set of slides](#) by Josh Tobin and this [article](#) regarding domain randomization for both vision and dynamics properties


2 Train your first RL agent

Soft deadline: 29/04/2026

Train a simple RL agent on the gym [Hopper](#) environment. This environment comes with an easy-to-use python interface which controls the underlying physics engine — [MuJoCo](#) — to model the robot.

The Hopper is a one-legged robot model whose task is to learn how to jump without falling, while achieving the highest possible horizontal speed.

Task 1 - The Gym Hopper environment



Check out the provided [code template](#) and start playing around with the underlying Hopper environment. Get familiar with the `test_random_policy.py` script, the python interface for

MuJoCo, the [Gym Documentation](#), and the Hopper environment overall. Finally answer the questions below.

Guiding Questions

- What is the state space in the Hopper environment? Is it discrete or continuous?
- What is the action space in the Hopper environment? Is it discrete or continuous?
- What is the mass value of each link of the Hopper environment, in the source and target variants respectively?

Hints

- If you need any help answering the above questions try looking at the [Mujoco documentation](#) or the [gym documentation](#).
- Bodies defined in the environment: `env.sim.model.body_names`
- Mass of all the corresponding bodies: `env.sim.model.body_mass`
- Number of degrees of freedom (DoFs) of the robot: `env.sim.model.nv`
- Number of DoFs for each body: `env.sim.model.body_dofnum`
- Number of actuators: `env.sim.model.nu`
- See other attributes [here](#)

2.1 Implement basic RL training algorithms

Implement from scratch two basic policy gradient reinforcement learning algorithms to train a simple control policy for the Hopper environment. Use `agent.py` for implementing the reinforcement learning algorithm itself (for example, the agent and policy classes). These classes are used to implement the main training loop in the `train.py` file. In particular, follow the tasks below, and make sure to go through the provided external resources:

Task 2 - REINFORCE (Vanilla Policy Gradient)

REINFORCE is a policy optimization algorithm that directly adjusts the parameters of a stochastic policy to maximize expected cumulative rewards and it operates through gradient ascent as optimization algorithm, iteratively searching for optimal parameters that maximize the objective function.

Background material:

- See Section 13.3-13.4 in [1] (if you want to get into theoretical details go through 13.-13.2 as well)
- See “REINFORCE” in [blog post](#)

Implement the following algorithms:

- REINFORCE without baseline,
- REINFORCE with a constant baseline $b = 20$

and compare their results.

Guiding Questions

- Analyze the performance of the trained policies in terms of reward and time consumption.
- How would you choose a good value for the baseline?

- How does the baseline affect the training, and why?

Task 3 - Actor Critic

Actor-Critic methods combine the advantages of value-based and policy gradient reinforcement learning.

Background material:

- See Section 13.5 in [1]
- See “Actor-Critic” in [blog post](#)

Implement the Actor-Critic Policy Gradient algorithm:

- refer to `train.py` as a starting point. It is okay to look at publicly available code for reference, but it's likely easier and more helpful to understand how to implement the code by yourself.


Guiding Questions

- Analyze the performance of the trained policies in terms of reward and time consumption.
- Compare the results with the REINFORCE algorithm you have previously obtained, highlighting any notable differences in terms of learning stability and convergence speed.

3 Lower/upper bound baselines

Soft deadline: 13/05/2025

PushTask of PandaGym



Push

The student will simulate the sim-to-real transfer scenario in a simplified sim-to-sim setting, as no work takes place on an actual real robot. In particular, two custom *domains* have been created ad-hoc: policy training takes place in the *source* environment and the student will transfer and test the policy on the *target* environment — which technically represents the real world. To simulate the reality gap, the source domain PushTask has been generated by shifting the cube mass by -4kg with respect to the target domain (where it is 5kg).

3.1 Implement advanced RL training pipelines

In this section, you'll make use of a third-party library to train an agent with state-of-the-art reinforcement learning algorithms such as PPO (Proximal Policy Optimization) and SAC (Soft Actor-Critic).

Stable-Baselines 3 is a state-of-the-art RL library offering efficient and stable implementations of various algorithms. Specifically, it provides robust implementations for PPO and SAC, making it an ideal choice for experimentation due to its ease of use and reliability.

Task 4 - Implement PPO and SAC

Follow the steps below, and make sure to go through the provided external resources:

1. Create a new script using the third-party library [stable-baselines3](#) (sb3) and train the *Push* agent with PPO [8] and SAC [7].
 - a. [openAI article on PPO](#)
 - b. [openAI article on SAC](#)
 - c. Explanation [video](#) on PPO, explanation [video](#) on SAC.
2. You may use the provided template in `train_sb3.py` as a starting point. It is okay to look at publicly available code for reference, but it's likely easier and more helpful to study the sb3 documentation and understand how to implement the code by yourself.
3. In the report you must report which works better and justify this result.

Hints:

- While PPO and SAC are more complex to understand, they will lead to better convergence during policy training and are likely less sensitive to hyperparameters.
- For the evaluation phase, refer to the "Evaluation Helper" in the [documentation](#).
- Feel free to adjust any hyperparameter such as learning rate or others as presented in the documentation if needed.
- You can use [callbacks for additional functionality](#), i.e., saving checkpoints, implementing early stopping, or integrating with TensorBoard or [WandB](#) for visualization.
- For the next section you will have to also implement a custom gymnasium [wrapper](#).

3.2 Training and testing

Task 5 - Train and test your policies

Train two agents with your algorithm of choice (the best performing), on the *source* and *target* domains respectively. Then, test each model and report its average return over 50 test episodes. In particular, report results for the following "training→test" configurations:

- source→source, (this will be just for reference, since the goal is to obtain optimal performance in the target environment)
- source→target (**lower bound**),
- target→target (**upper bound**).

Test with different hyperparameters and report the best results found together with the parameters used. The results will be the upper bound and lower bound for the following Domain Randomization phase.

Guiding Questions

- Why do we expect lower performances from the “source→target” configuration w.r.t. the “target→target”?
- If higher performances can be reached by training on the target environment directly, what prevents us from doing so (in a sim-to-real setting)?

4 Domain Randomization

Soft deadline: End of course

Implement Uniform Domain Randomization (UDR) and Automatic Domain Randomization (ADR) on the mass of the cube.

In this setting, UDR refers to manually designing a uniform distribution over a range of masses in the *source* environment (considering that the cube mass is fixed at -4kg w.r.t. the target one) and performing training with values that vary at each episode (sampled appropriately from the chosen distributions).

The underlying idea is to force the agent to maximize its reward and solve the task for a range of multiple environments at the same time, such that its learned behavior may be robust to dynamics variations.

Note that, since the choice of the distribution is a hyperparameter of the method, the student has to manually try different distribution ranges in order to expect good results on the target environment with UDR.

ADR is an extension of UDR that adaptively selects the ranges, please read the corresponding paper ([AutoDR](#)).

Task 6 - Implement Domain Randomization

Train a UDR and ADR agent on the *source* environment with the same RL algorithm previously used. Later test the policy obtained on both the *source* and *target* environments .

Guiding Questions

- Is UDR able to overcome the shift of mass and lead to more robust policies w.r.t. the naive “source→target” baseline in task 5?
- Can you think of limitations or downsides of UDR? What about ADR? Are there limitations/assumptions here?

Hints:

- *Remember not to act like you already know the target mass! Choosing a very narrow distribution for UDR around the target mass is basically cheating! We want a robust policy, not another mass-specific one!*

References

[1] “Reinforcement Learning: An introduction (Second Edition)” by Richard S. Sutton and Andrew G. Barto, [PDF](#)

[2] Kober, J., Bagnell, J. A., & Peters, J. (2013). “Reinforcement learning in robotics: A survey”. The International Journal of Robotics Research, [PDF](#)

[3] Kormushev, P., Calinon, S., & Caldwell, D. G. (2013). “Reinforcement learning in robotics: Applications and real-world challenges”, [PDF](#)

- [4] Höfer, S., Bekris, K., Handa, A., Gamboa, J. C., Golemo, F., Mozifian, M., ... & White, M. (2020). "Perspectives on sim2real transfer for robotics: A summary of the R: SS 2020 workshop", [PDF](#)
- [5] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World." arXiv, Mar. 20, 2017. [PDF](#)
- [6] Peng, X. B., Andrychowicz, M., Zaremba, W., & Abbeel, P. (2018, May). "Sim-to-real transfer of robotic control with dynamics randomization", [PDF](#)
- [7] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor.", [PDF](#)
- [8] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). "Proximal policy optimization algorithms", [PDF](#)
- [9] Muratore, Fabio, et al. "Robot learning from randomized simulations: A review." *Frontiers in Robotics and AI* 9 (2022)., [PDF](#)